

## JELILAT OLUWATOSIN ABDULLATEEF

Data Analyst

[abdullateefjelilat25@gmail.com](mailto:abdullateefjelilat25@gmail.com) (<mailto:abdullateefjelilat25@gmail.com>) /  
[www.linkedin.com/in/jelilat](http://www.linkedin.com/in/jelilat) (<http://www.linkedin.com/in/jelilat>)

## DIABETES PATIENT PREDICTION ANALYSIS

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether a patient has diabetes based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger dataset. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

### Importing Libraries

```
In [1]: #Importing Libraries:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#### *Importing libraries for prediction*

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
```

#### *loading the Dataset*

```
In [3]: df = pd.read_csv("diabetes.csv")
```

## Exploratory Data Analysis (EDA)

First five rows of the Dataset

In [4]: `df.head()`

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288

### Last five rows of the dataset

In [5]: `df.tail()`

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
763	10	101	76	48	180	32.9	0.17
764	2	122	70	27	0	36.8	0.34
765	5	121	72	23	112	26.2	0.24
766	1	126	60	0	0	30.1	0.34
767	1	93	70	31	0	30.4	0.31

### Overview of the dataset

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

### Shape of the dataset

```
In [7]: df.shape
```

```
Out[7]: (768, 9)
```

### Available Column names

```
In [8]: df.columns
```

```
Out[8]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
             dtype='object')
```

### Data type in each column

```
In [9]: df.dtypes
```

```
Out[9]: Pregnancies      int64  
         Glucose         int64  
         BloodPressure   int64  
         SkinThickness   int64  
         Insulin         int64  
         BMI             float64  
         DiabetesPedigreeFunction float64  
         Age            int64  
         Outcome        int64  
         dtype: object
```

## Checking for duplicate data

```
In [10]: df.duplicated().sum()
```

```
Out[10]: 0
```


## Checking for missing values

```
In [11]: df.isnull().sum()
```

```
Out[11]: Pregnancies      0  
         Glucose         0  
         BloodPressure   0  
         SkinThickness   0  
         Insulin         0  
         BMI             0  
         DiabetesPedigreeFunction 0  
         Age            0  
         Outcome        0  
         dtype: int64
```

```
In [12]: sns.heatmap(df.isnull())
```

```
Out[12]: <Axes: >
```



## Correlation Matrix

```
In [13]: correlation=df.corr()
print(correlation)
```

\	Pregnancies	Glucose	BloodPressure	SkinThickness
Pregnancies	1.000000	0.129459	0.141282	-0.081672
Glucose	0.129459	1.000000	0.152590	0.057328
BloodPressure	0.141282	0.152590	1.000000	0.207371
SkinThickness	-0.081672	0.057328	0.207371	1.000000
Insulin	-0.073535	0.331357	0.088933	0.436783
BMI	0.017683	0.221071	0.281805	0.392573
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928
Age	0.544341	0.263514	0.239528	-0.113970
Outcome	0.221898	0.466581	0.065068	0.074752


	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
In [14]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(),annot=True, cmap='coolwarm')
plt.title('Correlation', color = 'blue',fontweight='bold',fontsize=18)

plt.show
```

```
Out[14]: <function matplotlib.pyplot.show(close=None, block=None)>
```



## Training the Model with Train Test Split

### Train test split

Train-test split is a techniques used in machine learning to assess model performance. It divides the dataset into a training set and a testing set, with a 0.2 test size indicating that 20% of the data is used for testing and 80% for training.

```
In [15]: x=df.drop("Outcome",axis=1)
y=df['Outcome']
x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2, random_s
```

In X all the independent variables are stored In Y the predictor variable ("Outcome") is stored.

```
In [16]: scaler = StandardScaler()
x_train_scaler = scaler.fit_transform(x_train)
x_test_scaler = scaler.transform(x_test)

print(x_test_scaler.shape,x_train_scaler.shape)

(154, 8) (614, 8)
```

### Training the Model

Fitting the x train and y train data into the variable called model.

```
In [17]: model=LogisticRegression()
model.fit(x_train_scaler,y_train)
```

```
Out[17]: ▾ LogisticRegression
LogisticRegression()
```

## Making Prediction

```
In [18]: prediction = model.predict(x_test_scaler)
print(prediction)

[0 0 0 0 0 0 0 1 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 1 1 0 0 0 0 1 1 1 1 1 1 1
 0 1 1 0 1 1 0 0 1 1 0 0 1 0 1 1 0 0 0 1 0 0 1 1 0 0 0 0 1 0 1 0 1 1 0 0 0
 0 1 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 0 0 1 0 1 0 1 0 1 0 0 1 0 1 0
 0 0 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 1 1 1 1 0 0 1 0 0 1 1 0 0 0 0 0 0 0 0
 0 1 0 0 0 0]
```

After training the model, predictions are made using the data, which comprises 20% of the total datasets.

```
In [19]: accuracy = accuracy_score(prediction,y_test)
print('Accuracy:-',accuracy)
```

Accuracy:- 0.7532467532467533

## ACCURACY:- 75.32%

The model predicted the presence or absence of diabetes in approximately 75.32% of the cases

